

Hierarchical Visual Localization Using otgaussvld

Xiangyun He, Lin Ma, *Senior Member, IEEE*, Run Zhang, Yue Yu, Danyang Qin, Mu Zhou, *Senior Member, IEEE*

Abstract—Visual localization in large-scale scenes hinges on the quality of global retrieval. This paper presents otgaussvld, an optimal-transport-based Gaussian residual aggregator for place recognition. The proposed layer projects backbone tokens into a latent descriptor space, scores them with diagonal-Gaussian energy or normalized cosine similarity, and solves a Sinkhorn transport problem with a learnable dustbin and token-mass prediction. On top of the OT branch, we further introduce a self-attentive Gaussian query codebook and an OT-anchored refinement path that predicts bounded corrections for the image-side Gaussian moments without changing the descriptor dimensionality. Transport mass is retained at the cluster level and used to build mean and standard-deviation residuals with respect to learnable Gaussian prototypes. The concatenated 2KD-dimensional descriptor is directly ℓ_2 -normalized without an additional projection head. Experiments on Pittsburgh 30k and Aachen Day-Night demonstrate faster convergence and improved retrieval and localization performance over NetVLAD within a standard hierarchical localization pipeline.

Index Terms—Visual Localization, optimal transport, Gaussian mixture models, Sinkhorn iterations, place recognition.

I. INTRODUCTION

Visual localization typically involves 6 Degrees of Freedom (6DoF) positioning of a camera in three-dimensional space, which encompasses three degrees of freedom for three-dimensional translation and three degrees for three-dimensional rotation, corresponding to the three-dimensional position and three Euler angles in the positioning results. To provide positioning results within a consistent reference frame, 6DoF localization relies on a three-dimensional model of the scene, which can be obtained through Simultaneous Localization and Mapping (SLAM) or Structure from Motion (SfM). SLAM focuses more on the real-time aspect of mapping, discarding information and selecting only key frames for map construction during this process. SfM, on the other hand, places greater emphasis on the accuracy of the model, often at the expense of slower mapping speeds. Although both types of models are usable, in environments where real-time mapping requirements are less critical, there is a preference for the more accurate SfM models. In recent years,

visual localization methods based on three-dimensional scene models have primarily been categorized into two approaches: direct 2D-to-3D matching methods and image retrieval-based methods.

Direct 2D-to-3D matching methods involve comparing feature descriptors extracted from query images with those of 3D points in the scene's three-dimensional model to obtain 2D-3D match pairs [1], [2]. Since there may be some incorrect matches among these pairs, a Perspective-n-Point (PnP) algorithm based on Random Sample Consensus (RANSAC) is employed to estimate the camera pose from these 2D-3D matches, thereby filtering out the majority of incorrect matches. These methods typically rely on indexing techniques for matching, which can be computationally expensive, particularly in large-scale scenes with dense descriptor collections. Despite a series of improvements, there remains a trade-off between positioning accuracy and efficiency [3], [4].

Early image-based localization methods treated the localization problem as a place recognition issue, using the pose of the most similar image retrieved from a database as the pose of the query image. DenseVLAD [5] and NetVLAD [6] are representative of such methods. Both are effective in dealing with changes in lighting conditions and perform well in large-scale scenarios, although their positioning accuracy is somewhat lacking. Current image retrieval-based methods typically combine place recognition with indirect 2D-to-3D matching: initially, place recognition is used to narrow down the localization scene, followed by 2D-to-2D matching to establish 2D-to-3D correspondences, leading to refined pose estimation. These methods essentially employ a two-stage process for localization and are thus referred to as hierarchical localization based on image retrieval [7], [8].

Hierarchical localization based on image retrieval has demonstrated notable advantages in scalability, reliability, and effectiveness, establishing itself as a dominant paradigm for large-scale visual localization [8]. However, its performance critically depends on the quality of the global retrieval stage: failure to retrieve semantically coherent local regions leads to significant degradation in pose estimation accuracy. While increasing the number of retrieved images improves recall, it linearly escalates computational costs, hindering real-time applicability. In challenging scenarios such as low-texture areas, viewpoint changes, or illumination variations, the robustness of feature extraction is further compromised. Global retrieval methods like NetVLAD [6] rely solely on first-order statistical aggregation, lacking explicit modeling of local feature distribution differences, which frequently causes mismatches under cross-season or cross-view conditions. Meanwhile, local feature extractors face a trade-off between performance and efficiency: models such as SuperPoint [9] and DISK

Copyright (c) 20xx IEEE. Personal use of this material is permitted. However, permission to use this material for any other purposes must be obtained from the IEEE by sending a request to pubs-permissions@ieee.org.

This paper is supported by National Natural Science Foundation of China (62571164).

X. He, L. Ma, R. Zhang, Y. Yu are with the School of Electronics and Information Engineering, Harbin Institute of Technology, Harbin, China (e-mail: 24B305050@stu.hit.edu.cn; malin@hit.edu.cn; 25S105173@stu.hit.edu.cn; 24S105182@stu.hit.edu.cn). D. Qin is with the School of Electronics Engineering, Heilongjiang University, Harbin, China (e-mail: qindanyang@hlju.edu.cn). M. Zhou is with the School of Electronic Science and Engineering, Chongqing University of Posts and Telecommunications, Chongqing, China (e-mail: zhoumu@cqupt.edu.cn). (Corresponding author: Lin Ma.)

[10] achieve high discriminability at the cost of substantial computational overhead, while lightweight alternatives like XFeat [11] suffer from insufficient salient region perception due to overly simplified network architectures, limiting feature distinctiveness.

To effectively address the aforementioned issues, we have made the following contributions:

(1) We propose otgaussvld, a global feature aggregation framework that couples diagonal-Gaussian scoring with Sinkhorn transport and a dustbin, thereby producing assignments under explicit mass constraints rather than independent token-wise softmax normalization.

(2) We design a self-attentive Gaussian query codebook and an OT-anchored refinement branch, which use bounded residual gates to adapt prototype parameters and image-side Gaussian moments while remaining close to the base OT solution at initialization.

(3) We derive a transport-aware Gaussian residual descriptor that preserves cluster mass, models both mean and scale deviations, and admits a diagonal-Gaussian Wasserstein interpretation of the resulting cluster residuals.

(4) Extensive experiments show that the proposed aggregator improves retrieval accuracy and training stability over NetVLAD while remaining fully compatible with standard hierarchical localization back ends.

II. PROPOSED METHOD

A. Framework

The hierarchical visual localization framework consists of two primary stages: an offline mapping phase and an online localization phase, with the overall workflow illustrated in Fig. 1. The objective of this framework is to efficiently and robustly estimate the 3D position and orientation (6-DoF pose) of a camera based on a single query image.

The core objective of the offline mapping phase is to construct a sparse 3D point cloud model (map) of the target scene and precompute essential features. Specifically, we employ an incremental Structure-from-Motion (SfM) algorithm for 3D reconstruction. This process begins by processing a collection of database images using a sparse local feature extractor and the proposed otgaussvld global descriptor. The extracted local features are used for feature matching and triangulation steps within SfM to recover the 3D structure and the precise pose for each database image. The global descriptors are also computed and stored. The final mapping results, which encompass the 3D point cloud, the database images' poses, their corresponding local features, and global descriptors, are stored in a database to support efficient online localization.

The online localization phase employs a coarse-to-fine pipeline, comprising three sequential steps. First, in the global retrieval stage, the query image is processed by the otgaussvld method to extract a global feature descriptor. This descriptor enables rapid similarity computation and K-Nearest Neighbor (KNN) search against precomputed database features, yielding a shortlist of visually similar candidate images. Second, during local feature matching and 2D-3D correspondence establishment, the query and top candidate images are processed by

a sparse local feature extractor to obtain local descriptors. Robust 2D-2D feature matching is first established between the query image and candidate database images. These matches are subsequently converted into 2D-3D point correspondences for the query by leveraging the pre-existing associations between 2D keypoints in the candidate images and 3D points in the Structure-from-Motion (SfM) model. Finally, for pose estimation, a RANSAC-based Perspective-n-Point algorithm utilizes the 2D-3D correspondences to estimate the final 6-DoF camera pose.

B. Transport-Aware Gaussian Residual Aggregator

Let $X \in \mathbb{R}^{C \times H \times W}$ denote the backbone output. After the two-layer point-wise adapter and LayerNorm, the feature map is flattened into $N = HW$ local tokens $z_n \in \mathbb{R}^D$. In the default configuration, $C = 768$, $D = 64$, and $K = 64$, resulting in a raw descriptor dimension of $2KD = 8192$. No projection head is applied after aggregation.

For each cluster $k \in \{1, \dots, K\}$, the layer parameterizes a diagonal Gaussian prototype

$$p_k(z) = \mathcal{N}(z; \mu_k, \text{diag}(\sigma_k^2)),$$

where $\mu_k \in \mathbb{R}^D$, $\sigma_k \in \mathbb{R}_+^D$, and the mixture prior is

$$\pi_k = \frac{e^{\alpha_k}}{\sum_{j=1}^K e^{\alpha_j}}.$$

The default GMM score is the corresponding log-density up to a cluster-independent constant:

$$s_{kn} = \log \pi_k - \frac{1}{2} \sum_{d=1}^D \frac{(z_{nd} - \mu_{kd})^2}{\sigma_{kd}^2} - \sum_{d=1}^D \log \sigma_{kd}. \quad (1)$$

For completeness, the attention ablation uses normalized prototypes and tokens, $\bar{\mu}_k = \mu_k / \|\mu_k\|_2$ and $\bar{z}_n = z_n / \|z_n\|_2$, with score $s_{kn}^{\text{att}} = \tau \langle \bar{\mu}_k, \bar{z}_n \rangle$, where $\tau = \text{softplus}(a) + \epsilon$.

When `token_mass_mode` is set to `learned`, a lightweight token saliency head produces logits q_n and masses

$$\nu_n = \frac{\exp(\eta q_n)}{\sum_{m=1}^N \exp(\eta q_m)}, \quad \eta = \text{softplus}(\lambda) + \epsilon.$$

In the uniform case, $\nu_n = 1/N$. The target mass vector is $n = \nu / \|\nu\|_1$.

To suppress uninformative tokens, the score matrix is augmented with a learnable dustbin row and the following entropy-regularized transport problem is solved:

$$A^* = \arg \max_{A \in \mathcal{U}(m, n)} \langle \tilde{S}, A \rangle + \mathcal{H}(A),$$

where $\mathcal{U}(m, n) = \{A \geq 0 : A\mathbf{1} = m, A^\top \mathbf{1} = n\}$, $\mathcal{H}(A) = -\sum_{k,n} A_{kn} \log A_{kn}$, and

$$m = [(1 - \delta)\pi_1, \dots, (1 - \delta)\pi_K, \delta]^\top.$$

Here $\delta \in (0, 1)$ is the dustbin mass and $\tilde{S}$ denotes the score matrix after appending the dustbin row. The implemented log-space Sinkhorn iterations compute the associated dual scalings.

Proposition 1 (Well-posedness of the transport assignment). *If the source mass m and target mass n are strictly positive,*




Fig. 1. Flowchart of the hierarchical visual localization system. The system is divided into two phases: offline and online. In the offline phase, database images are processed by a sparse local feature extractor and otgaussvlad to construct a database containing a global index and a sparse 3D point cloud (SfM Reconstruction). In the online phase, a query image is first processed by otgaussvlad for global retrieval to identify candidate scene regions. Subsequently, sparse local features are matched with the retrieved candidates in a 2D-3D matching step. Finally, the precise 6-DoF camera pose is estimated using a PnP algorithm.

then the entropy-regularized transport problem admits at least one maximizer. Because the entropy term is strictly concave on the interior of $\mathcal{U}(m, n)$, the maximizer is unique.

Proof. The feasible set $\mathcal{U}(m, n)$ is non-empty, closed, and bounded, hence compact. The map $A \mapsto \langle \tilde{S}, A \rangle$ is linear, and $\mathcal{H}(A)$ is strictly concave on the positive orthant. Therefore the objective is strictly concave over a compact convex set, which yields existence and uniqueness of the maximizer. $\square$

Let A be the real-cluster block of A^* . For each cluster k , define the transported mass

$$m_k = \sum_{n=1}^N A_{kn}, \quad \gamma_{kn} = \frac{A_{kn}}{\max(m_k, \epsilon)}.$$

If $m_k > 0$, then $\sum_n \gamma_{kn} = 1$, so γ_k defines a probability distribution over the tokens assigned to cluster k . The transport-induced empirical moments are then

$$\hat{\mu}_k = \sum_{n=1}^N \gamma_{kn} z_n,$$

$$\hat{v}_k = \sum_{n=1}^N \gamma_{kn} (z_n - \hat{\mu}_k) \odot (z_n - \hat{\mu}_k), \quad (3)$$

$$\hat{\sigma}_k = \sqrt{\hat{v}_k}. \quad (4)$$

Proposition 2 (Transport-induced Gaussian moments). *The quantities $\hat{\mu}_k$ and $\hat{v}_k$ are exactly the first and diagonal second central moments of the empirical distribution $q_k = \sum_n \gamma_{kn} \delta_{z_n}$. Equivalently, $\hat{\mu}_k$ is the unique minimizer of $\sum_n \gamma_{kn} \|z_n - \mu\|_2^2$.*

Proof. Because $\sum_n \gamma_{kn} = 1$, q_k is a valid discrete probability measure. Its mean is

$$\mathbb{E}_{q_k}[z] = \sum_{n=1}^N \gamma_{kn} z_n = \hat{\mu}_k,$$

which establishes the first identity. To obtain the minimizer characterization, define the weighted quadratic objective

$$\phi_k(\mu) = \sum_{n=1}^N \gamma_{kn} \|z_n - \mu\|_2^2.$$

Expanding the square gives

$$\phi_k(\mu) = \sum_{n=1}^N \gamma_{kn} \|z_n\|_2^2 - 2\mu^\top \sum_{n=1}^N \gamma_{kn} z_n + \|\mu\|_2^2. \quad (2)$$

Hence

$$\nabla_\mu \phi_k(\mu) = 2\left(\mu - \sum_{n=1}^N \gamma_{kn} z_n\right), \quad \nabla_\mu^2 \phi_k(\mu) = 2I \succ 0, \quad (4)$$

so ϕ_k is strictly convex and has a unique minimizer at $\mu = \hat{\mu}_k$.

For the diagonal second moment, expand the centered square elementwise:

$$\begin{aligned}\hat{v}_k &= \sum_{n=1}^N \gamma_{kn} (z_n - \hat{\mu}_k) \odot (z_n - \hat{\mu}_k) \\ &= \sum_{n=1}^N \gamma_{kn} (z_n \odot z_n) - \hat{\mu}_k \odot \hat{\mu}_k.\end{aligned}\quad (5)$$

This is precisely the elementwise variance of q_k . Therefore $\hat{v}_k$ and $\hat{\sigma}_k$ are the diagonal covariance and standard deviation of the transported component. $\square$

The cluster residual is formed as

$$r_k = [\hat{\mu}_k - \mu_k, \hat{\sigma}_k - \sigma_k].$$

This residual admits a natural interpretation in the diagonal-Gaussian Wasserstein geometry. In particular, once the covariance is expressed through its standard deviation coordinates, the 2-Wasserstein metric becomes Euclidean. The concatenation of cluster-wise residuals therefore yields the final global descriptor.

Proposition 3 (Wasserstein feature representation). *Using the closed-form 2-Wasserstein distance between Gaussian measures [24], for each Gaussian component the squared 2-Wasserstein distance between the current image component $X_i \sim \mathcal{N}(\hat{\mu}_i, \hat{\Sigma}_i)$ and the learned prototype component $Y_i \sim \mathcal{N}(\mu_i, \Sigma_i)$ admits the closed form*

$$\begin{aligned}W_2^2(X_i, Y_i) &= \|\hat{\mu}_i - \mu_i\|_2^2 \\ &\quad + \text{Tr} \left[\hat{\Sigma}_i + \Sigma_i - 2 \left(\hat{\Sigma}_i^{1/2} \Sigma_i \hat{\Sigma}_i^{1/2} \right)^{1/2} \right].\end{aligned}\quad (6)$$

When the covariance matrices are diagonal, write $\hat{\Sigma}_i = \text{diag}(\hat{\sigma}_i^2)$ and $\Sigma_i = \text{diag}(\sigma_i^2)$. Since diagonal matrices commute,

$$\left(\hat{\Sigma}_i^{1/2} \Sigma_i \hat{\Sigma}_i^{1/2} \right)^{1/2} = \text{diag}(\hat{\sigma}_i \odot \sigma_i). \quad (7)$$

and the trace term reduces to

$$\text{Tr} \left[\hat{\Sigma}_i + \Sigma_i - 2 \left(\hat{\Sigma}_i^{1/2} \Sigma_i \hat{\Sigma}_i^{1/2} \right)^{1/2} \right] = \|\hat{\sigma}_i - \sigma_i\|_2^2. \quad (8)$$

We therefore define the component-wise Wasserstein feature as

$$g_i = [\hat{\mu}_i - \mu_i, \text{flatten}(\hat{\Sigma}_i^{1/2} - \Sigma_i^{1/2})].$$

For the diagonal case, $\text{flatten}(\hat{\Sigma}_i^{1/2} - \Sigma_i^{1/2})$ is equivalent to $\hat{\sigma}_i - \sigma_i$. In this coordinate system, the feature map is an isometric embedding of the diagonal-Gaussian Wasserstein geometry into Euclidean space.

Proof. For the first claim, substitute the diagonal parameterization $\hat{\Sigma}_i = \text{diag}(\hat{\sigma}_i^2)$ and $\Sigma_i = \text{diag}(\sigma_i^2)$. Because diagonal matrices commute, the covariance term simplifies to

$$\begin{aligned}\text{Tr} \left(\hat{\Sigma}_i + \Sigma_i - 2 \left(\hat{\Sigma}_i^{1/2} \Sigma_i \hat{\Sigma}_i^{1/2} \right)^{1/2} \right) &= \sum_{d=1}^D (\hat{\sigma}_{id} - \sigma_{id})^2 \\ &= \|\hat{\sigma}_i - \sigma_i\|_2^2.\end{aligned}\quad (9)$$

Adding the mean term yields

$$\begin{aligned}\|g_i\|_2^2 &= \|\hat{\mu}_i - \mu_i\|_2^2 + \|\hat{\sigma}_i - \sigma_i\|_2^2 \\ &= W_2^2(X_i, Y_i),\end{aligned}\quad (10)$$

which proves the first claim.

For the second claim, the prototype component is shared, so the prototype offsets cancel when the two Wasserstein features are subtracted:

$$g_i - g'_i = [\hat{\mu}_i - \hat{\mu}'_i, \text{flatten}(\hat{\Sigma}_i^{1/2} - \hat{\Sigma}'_i^{1/2})].$$

Applying the same diagonal reduction to the pair $(\hat{\mu}_i, \hat{\sigma}_i)$ and $(\hat{\mu}'_i, \hat{\sigma}'_i)$ yields

$$\begin{aligned}\|g_i - g'_i\|_2^2 &= \|\hat{\mu}_i - \hat{\mu}'_i\|_2^2 + \|\hat{\sigma}_i - \hat{\sigma}'_i\|_2^2 \\ &= \|\hat{\mu}_i - \hat{\mu}'_i, \hat{\sigma}_i - \hat{\sigma}'_i\|_2^2,\end{aligned}\quad (11)$$

which is exactly the squared 2-Wasserstein distance between the two diagonal Gaussian components. Consequently, concatenating the cluster-wise residuals produces a vector-valued embedding of component-wise Wasserstein discrepancies rather than a scalar distance. $\square$

Remark. In the diagonal parameterization used here, the second coordinate is the standard deviation σ , equivalently the square root of the diagonal variance; this is precisely the coordinate system in which the diagonal-Gaussian 2-Wasserstein metric becomes Euclidean.

When `mass_preserving` is enabled, the residual is reweighted by

$$w_k = \left(\frac{K}{1 - \delta} m_k \right)^\beta, \quad \beta = \text{mass_power},$$

so that the final cluster residual becomes $\tilde{r}_k = w_k r_k$. This normalization is anchored at the uniform reference mass: if $m_k = (1 - \delta)/K$, then $w_k = 1$. The final global descriptor is obtained by concatenating $\tilde{r}_1, \dots, \tilde{r}_K$ and applying ℓ_2 normalization:

$$g = \frac{\text{vec}([\tilde{r}_1, \dots, \tilde{r}_K])}{\|\text{vec}([\tilde{r}_1, \dots, \tilde{r}_K])\|_2}.$$

Therefore, the layer can be viewed as a prototype-indexed entropy-regularized transport operator followed by a Wasserstein residual encoder.

C. Self-Attentive Gaussian Query Codebook

The implementation in `ot_gaussvld_gaussian_query_codebook` makes the Gaussian codebook slightly adaptive by replacing the fixed prototype table with a learnable query set. Let $u_k \in \mathbb{R}^{C_g}$ denote the k -th Gaussian query slot and $Q(\cdot)$ be a self-attentive block with optional residual attention and LayerNorm. The slot representation is

$$h_k = Q(u_k), \quad (12)$$

followed by two bounded residual heads:

$$\Delta \mu_k = s_\mu \tanh(W_\mu h_k), \quad (13)$$

$$\Delta \ell_k = s_\sigma \tanh(W_\ell h_k), \quad (14)$$

where $\ell_k = \log \sigma_k$, s_μ , and s_σ are learnable scales initialized to small positive values through an inverse-sigmoid parameterization. The adaptive Gaussian codebook is then

$$\tilde{\mu}_k = \mu_k + \Delta\mu_k, \quad \tilde{\ell}_k = \ell_k + \Delta\ell_k. \quad (15)$$

Because the residual heads are zero-initialized and the scales start small, the module is close to the original prototype table at the beginning of training and can only deviate when the data supports it.

D. OT-Anchored Moment Refinement

The stronger variant in `ot_gaussvld_gaussian_query_codebook_ot_anchored_refine.py` keeps the OT assignment branch intact, but refines the image-side Gaussian statistics after transport. Given the transported assignments A , the OT branch first produces

$$\hat{\mu}_k^{\text{ot}} = \sum_n \gamma_{kn} z_n, \quad \hat{\sigma}_k^{\text{ot}} = \sqrt{\sum_n \gamma_{kn} z_n^2 - (\hat{\mu}_k^{\text{ot}})^2}, \quad (16)$$

where γ_{kn} is the row-normalized assignment.

The refiner is parameterized as a small transformer-style block. Its input query is the concatenation $[\mu_k, \log \sigma_k]$, which is first normalized by a 2D-dimensional LayerNorm and then projected by a two-layer MLP $\text{Linear}(2D \rightarrow C) \rightarrow \text{GELU} \rightarrow \text{Dropout} \rightarrow \text{Linear}(C \rightarrow C)$. This produces a cluster query in the latent refinement space. The local tokens are similarly normalized by a D -dimensional LayerNorm and projected with a linear map into the same latent dimension so that query and memory share the attention embedding space.

The core interaction is a batch-first multi-head cross-attention module,

$$\text{Attn}(Q, M) = \text{MHA}(Q, M, M), \quad (17)$$

where the query sequence is the Gaussian codebook and the key/value sequence is the projected local-token memory. The attention output is added back to the query branch through a residual connection, followed by LayerNorm. If enabled, a position-wise feed-forward block is applied on top of the attention features:

$$\text{FFN}(x) = W_2 \text{Dropout}(\text{GELU}(W_1 x)), \quad (18)$$

with residual addition after a second normalization. This makes the refiner a lightweight transformer encoder rather than a single attention layer.

From the refined latent features, two independent linear heads predict the residual corrections for mean and log-standard-deviation:

$$\Delta\mu_k = W_\mu h_k + b_\mu, \quad (19)$$

$$\Delta\ell_k = W_\ell h_k + b_\ell. \quad (20)$$

Both heads are explicitly zero-initialized, so the refiner begins at the identity mapping and initially leaves the OT moments unchanged. Optionally, a tanh nonlinearity is applied to bound the raw residual before scaling. The residual magnitude is then multiplied by learnable gates

$$\alpha_\mu = s_{\max} \sigma(\rho_\mu), \quad \alpha_\sigma = s_{\max} \sigma(\rho_\sigma), \quad (21)$$

where $s_{\max}$ is the maximum residual scale and $\sigma(\cdot)$ is the sigmoid function. The logit parameters ρ_μ and ρ_σ are initialized using an inverse-sigmoid map so that the effective scales start at small positive values (e.g., 0.05 by default). This guarantees that the correction branch is conservative at the beginning of training and can only deviate gradually from the OT estimate.

The final moment estimates are anchored to the OT statistics rather than replacing them:

$$\hat{\mu}_k = \hat{\mu}_k^{\text{ot}} + \Delta\mu_k, \quad \hat{\ell}_k = \log(\hat{\sigma}_k^{\text{ot}}) + \Delta\ell_k. \quad (22)$$

Finally, the image-side scale is recovered by exponentiation and clamped to a small positive constant for numerical stability, matching the implementation of $\hat{\sigma}_k = \exp(\hat{\ell}_k)$. This design is important in practice: the refinement branch starts almost inactive, so optimization initially follows the stable OT solution and then learns image-side corrections only when they reduce retrieval loss.

III. EXPERIMENTAL RESULTS AND ANALYSIS

A. Experiment Setup

This section details the implementation specifics and experimental configurations for the proposed `otgaussvld` method. The setup encompasses training datasets, parameter configurations, and hardware/software environments to ensure reproducibility.

1) *otgaussvld*: The training dataset, Pittsburgh 30k [6], comprises approximately 30,000 geotagged street-view images from Google Street View, with GPS coordinates accurate within 7–15 meters. Diversity in capture times, viewing angles, and lighting conditions offers a wide range of positive and hard negative samples. Experiments were conducted on a system equipped with an NVIDIA GeForce RTX 4080 GPU and an Intel Core i9-14900K CPU. The software environment included CUDA 12.4, NVIDIA driver version 550.54.14, PyTorch 1.13.0 with CUDA 11.7 support, and torchvision 0.14.0. In the `otgaussvld` configuration used in our experiments, the aggregator is instantiated with 768 input channels, 64 clusters, a 64-dimensional cluster space, 768 hidden channels, GMM scoring, three Sinkhorn iterations, a dustbin mass of 0.05, learned token mass prediction, mass-preserving residuals, and `mass_power` set to 1.0. This configuration yields an 8192-dimensional global descriptor. The training hyperparameters for both NetVLAD and the proposed `otgaussvld` method were identical: an L2 regularization coefficient of 0.001, initial learning rate of 0.0001, learning rate decay rate of 0.5 applied every 5 epochs, momentum of 0.9, and triplet margin threshold of 0.3.

B. *otgaussvld*

Metrics. Recall@ K is a core metric for global image retrieval, measuring the proportion of correctly matched images among the top- K database results. It evaluates feature encoding efficacy: high Recall@1 reflects top-tier precision, while Recall@5 and higher thresholds indicate robustness to challenges like illumination or viewpoint variations. This ensures reliable candidate image selection for subsequent local matching and camera pose estimation.

TABLE I
COMPARISON OF RECALL RATES ON THE TEST SET

Method	Pittsburgh 250k-test			
	Recall@1	Recall@5	Recall@10	Recall@20
SPE-NetVLAD [14]	85.60%	93.20%	94.55%	95.90%
NetVLAD [6]	81.64%	90.98%	93.26%	95.05%
MR-NetVLAD [17]	86.70%	93.60%	94.80%	96.00%
Ours(VGG)	84.42%	92.74%	94.87%	96.20%
Ours(EfficientNet)	88.42%	95.91%	97.39%	98.31%

Results. Fig. 2 demonstrates otgaussvlad’s significant performance advantage over NetVLAD. Initially, otgaussvlad achieves a lower loss (0.05 vs. NetVLAD’s 0.23), indicating superior initialization. It converges rapidly, stabilizing near 0.01 after merely 5 epochs and approaching zero by epoch 15. In contrast, NetVLAD exhibits slow convergence, with loss remaining at 0.17 even after 15 epochs. The results confirm otgaussvlad’s excellence in training efficiency, model performance, and stability.




Fig. 2. Comparison of training loss over epochs. The graph shows the change in loss during the training process for NetVLAD and otgaussvlad methods. The x-axis represents the number of training epochs, and the y-axis represents the training loss.

The comparative evaluation of recall rates on the Pittsburgh 250k-test dataset is summarized in Table I. The proposed method with an EfficientNet backbone achieves state-of-the-art performance across all evaluation metrics, attaining recall rates of 88.42% (Recall@1), 95.91% (Recall@5), 97.39% (Recall@10), and 98.31% (Recall@20). This represents a significant improvement over existing approaches, outperforming the second-best method (MR-NetVLAD) by 1.72% in Recall@1 and by 2.31% in Recall@5.

The comparison highlights the critical role of backbone architecture selection: the EfficientNet variant exhibits a substantial gain over the VGG-based implementation, particularly in Recall@1 (+4.00%) and Recall@5 (+3.17%). Both variants consistently surpass the NetVLAD baseline, validating the efficacy of the proposed feature encoding framework in large-scale visual localization tasks.

To eliminate the impact of the training framework and backbone network on model performance, we uniformly used the VGG network as the Backbone and retrained NetVLAD and otgaussvlad using PyTorch. To verify the effectiveness of our proposed global retrieval method in localization, we evaluated the role of our method otgaussvlad in localization using the hierarchical localization toolbox (Hloc) under fixed

sparse local features and feature matching methods. As shown in Table II, the results indicate that our method can recall database images highly relevant to the query images for both daytime and nighttime scenes. This outcome aligns perfectly with the essence of our design for otgaussvlad, which employs the Wasserstein distance as a measure of feature similarity. This approach better considers the distribution of local semantic features, giving the extracted global features a significant advantage when retrieving candidate images.

TABLE II
PERFORMANCE COMPARISON OF GLOBAL FEATURE EXTRACTING METHODS ON AACHEN DAY-NIGHT

Method	Day			Night		
	0.25m 2°	0.5m 5°	5m 10°	0.25m 2°	0.5m 5°	5m 10°
NetVLAD [6]	81.6	87.4	92.0	71.4	81.6	91.8
otgaussvlad	83.1	89.9	93.9	79.6	87.8	94.9

IV. CONCLUSION

This paper presents otgaussvlad, a transport-aware global feature aggregator for hierarchical visual localization. By combining diagonal-Gaussian scoring, Sinkhorn transport with a dustbin, learned token mass, and mass-preserving mean/scale residuals, the proposed layer produces a compact descriptor that is more robust to appearance variation than NetVLAD while remaining fully compatible with conventional retrieval-and-matching pipelines. Experiments on Pittsburgh 30k and Aachen Day-Night confirm faster convergence and improved retrieval and localization performance.

REFERENCES

- [1] K. Hajebi, Y. Abbasi-Yadkori, H. Shahbazi, and H. Zhang, “Fast approximate nearest-neighbor search with k-nearest neighbor graph,” in *Proceedings of the Twenty-Second International Joint Conference on Artificial Intelligence - Volume Volume Two*, ser. IJCAI’11. AAAI Press, 2011, p. 1312–1317.
- [2] Y. Li, N. Snavey, and D. P. Huttenlocher, “Location recognition using prioritized feature matching,” in *Proceedings of the 11th European Conference on Computer Vision: Part II*, ser. ECCV’10. Berlin, Heidelberg: Springer-Verlag, 2010, p. 791–804.
- [3] T. Sattler, B. Leibe, and L. Kobbelt, “Efficient & effective prioritized matching for large-scale image-based localization,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 39, no. 9, pp. 1744–1756, 2017.
- [4] Z. Song, C. Wang, Y. Liu, and S. Shen, “Recalling direct 2d-3d matches for large-scale visual localization,” in *2021 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2021, pp. 1191–1197.
- [5] R. Arandjelovic and A. Zisserman, “All about vlad,” in *2013 IEEE Conference on Computer Vision and Pattern Recognition*, 2013, pp. 1578–1585.
- [6] R. Arandjelović, P. Gronat, A. Torii, T. Pajdla, and J. Sivic, “Netvlad: Cnn architecture for weakly supervised place recognition,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 40, no. 6, pp. 1437–1451, 2018.
- [7] T. Shi, S. Shen, X. Gao, and L. Zhu, “Visual localization using sparse semantic 3d map,” *2019 IEEE International Conference on Image Processing (ICIP)*, pp. 315–319, 2019.
- [8] P.-E. Sarlin, C. Cadena, R. Siegwart, and M. Dymczyk, “From coarse to fine: Robust hierarchical localization at large scale,” in *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 12 708–12 717.
- [9] D. Detone, T. Malisiewicz, and A. Rabinovich, “Superpoint: Self-supervised interest point detection and description,” 2017.

- [10] M. J. Tyszkiewicz, P. Fua, and E. Trulls, "Disk: Learning local features with policy gradient," 2020.
- [11] G. Potje, F. Cadar, A. Araujo, R. Martins, and E. R. Nascimento, "Xfeat: Accelerated features for lightweight image matching," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024, pp. 2682–2691.
- [12] Y. Zhong, R. Arandjelović, and A. Zisserman, "Ghostvlad for set-based face recognition," in *Asian conference on computer vision*. Springer, 2018, pp. 35–50.
- [13] A. Khaliq, S. Ehsan, Z. Chen, M. Milford, and K. McDonald-Maier, "A holistic visual place recognition approach using lightweight cnns for significant viewpoint and appearance changes," *IEEE transactions on robotics*, vol. 36, no. 2, pp. 561–569, 2019.
- [14] J. Yu, C. Zhu, J. Zhang, Q. Huang, and D. Tao, "Spatial pyramid-enhanced netvlad with weighted triplet loss for place recognition," *IEEE transactions on neural networks and learning systems*, vol. 31, no. 2, pp. 661–674, 2019.
- [15] Y. Ge, H. Wang, F. Zhu, R. Zhao, and H. Li, "Self-supervising fine-grained region similarities for large-scale image localization," in *European conference on computer vision*. Springer, 2020, pp. 369–386.
- [16] S. Hausler, S. Garg, M. Xu, M. Milford, and T. Fischer, "Patch-netvlad: Multi-scale fusion of locally-global descriptors for place recognition," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2021, pp. 14 141–14 152.
- [17] A. Khaliq, M. Milford, and S. Garg, "Multires-netvlad: Augmenting place recognition training with low-resolution imagery," *IEEE Robotics and Automation Letters*, vol. 7, no. 2, pp. 3882–3889, 2022.
- [18] Z. Li, C. D. W. Lee, B. X. L. Tung, Z. Huang, D. Rus, and M. H. Ang, "Hot-netvlad: Learning discriminatory key points for visual place recognition," *IEEE Robotics and Automation Letters*, vol. 8, no. 2, pp. 974–980, 2023.
- [19] E. Rublee, V. Rabaud, K. Konolige, and G. Bradski, "Orb: An efficient alternative to sift or surf," in *2011 International conference on computer vision*. Ieee, 2011, pp. 2564–2571.
- [20] P. Gleize, W. Wang, and M. Feiszli, "Silk: Simple learned keypoints," in *Proceedings of the IEEE/CVF international conference on computer vision*, 2023, pp. 22 499–22 508.
- [21] X. Zhao, X. Wu, J. Miao, W. Chen, P. C. Y. Chen, and Z. Li, "Alike: Accurate and lightweight keypoint detection and descriptor extraction," *IEEE Transactions on Multimedia*, vol. 25, pp. 3101–3112, 2023.
- [22] M. Kanakis, S. Maurer, M. Spallanzani, A. Chhatkuli, and L. Van Gool, "Zippypoint: Fast interest point detection, description, and matching through mixed precision discretization," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023, pp. 6114–6123.
- [23] M. Tan and Q. V. Le, "Efficientnet: Rethinking model scaling for convolutional neural networks," 2019.
- [24] I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, and A. C. Courville, "Improved training of wasserstein gans," *Advances in neural information processing systems*, vol. 30, 2017.
- [25] A. G. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam, "Mobilenets: Efficient convolutional neural networks for mobile vision applications," 2017.
- [26] D. Yin, L. Hu, B. Li, Y. Zhang, and X. Yang, "5% $\rightarrow$ 100%: Breaking performance shackles of full fine-tuning on visual recognition tasks," in *Proceedings of the Computer Vision and Pattern Recognition Conference*, 2025, pp. 20 071–20 081.
- [27] S. Woo, J. Park, J.-Y. Lee, and I. S. Kweon, "Cbam: Convolutional block attention module," in *Proceedings of the European conference on computer vision (ECCV)*, 2018, pp. 3–19.
- [28] J. Hu, L. Shen, and G. Sun, "Squeeze-and-excitation networks," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2018, pp. 7132–7141.
- [29] Q. Wang, B. Wu, P. Zhu, P. Li, W. Zuo, and Q. Hu, "Eca-net: Efficient channel attention for deep convolutional neural networks," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2020, pp. 11 534–11 542.
- [30] J. Chen, S.-h. Kao, H. He, W. Zhuo, S. Wen, C.-H. Lee, and S.-H. G. Chan, "Run, don't walk: chasing higher flops for faster neural networks," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2023, pp. 12 021–12 031.
- [31] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga *et al.*, "Pytorch: An imperative style, high-performance deep learning library," *Advances in neural information processing systems*, vol. 32, 2019.
- [32] Z. Li and N. Snavely, "Megadepth: Learning single-view depth prediction from internet photos," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2018, pp. 2041–2050.
- [33] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, "Microsoft coco: Common objects in context," in *European conference on computer vision*. Springer, 2014, pp. 740–755.
- [34] A. Dai, A. X. Chang, M. Savva, M. Halber, T. Funkhouser, and M. Nießner, "ScanNet: Richly-annotated 3d reconstructions of indoor scenes," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2017, pp. 5828–5839.
- [35] P. Lindenberger, P.-E. Sarlin, and M. Pollefeys, "Lightglue: Local feature matching at light speed," in *Proceedings of the IEEE/CVF international conference on computer vision*, 2023, pp. 17 627–17 638.
- [36] J. Sun, Z. Shen, Y. Wang, H. Bao, and X. Zhou, "Loftr: Detector-free local feature matching with transformers," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2021, pp. 8922–8931.
- [37] V. Balntas, K. Lenc, A. Vedaldi, and K. Mikolajczyk, "Hpatches: A benchmark and evaluation of handcrafted and learned local descriptors," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2017, pp. 5173–5182.
- [38] T. Sattler, W. Maddern, C. Toft, A. Torii, L. Hammarstrand, E. Stenborg, D. Safari, M. Okutomi, M. Pollefeys, J. Sivic *et al.*, "Benchmarking 6dof outdoor visual localization in changing conditions," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2018, pp. 8601–8610.
- [39] J. L. Schonberger and J.-M. Frahm, "Structure-from-motion revisited," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2016, pp. 4104–4113.